Kehang Zhu

Contact

kehangzhu (at) stanford.edu ◇ kehang-zhu.github.io

Education

Stanford University, Digital Economy Lab 2026 - Present
Postdoctoral Scholar, Stanford Institute for Human-Centered Artificial Intelligence (HAI)

Harvard University
Ph.D. Physics 09/2021 - 06/2026

Harvard University
S.M., Computer Science 05/2024

University of Science and Technology of China
B.A. Physics (*summa cum laude*) 05/2021

References

John J. Horton (Chair)
MIT Sloan School of Management
jjhorton (at) mit.edu

David C. Parkes
Harvard University
parkes (at) eecs.harvard.edu

Nithum Thain
Google DeepMind
nthain (at) google.com

Erik Brynjolfsson
Stanford University
erikb (at) stanford.edu

Grants ($150,000)

Google DeepMind Seed Fund ($50,000; Co-PI w/ John J. Horton) 2024
DropBox Research Fund ($100,000; Co-PI w/ John J. Horton) 2023

Awards & Scholarships

Harvard Student Teaching Award 2024
Introduction to Technical AI Safety Fellowship 2023
Harvard Purcell Fellowship 2021

Working Papers

Designing Behavioral Profiles for Language Model Simulation of Out-of-Distribution Human Behavior
(Robin Na, **Kehang Zhu**, Logan Cross, & Michiel A. Bakker)
Workshop on User Simulation (UserSim) at *NeurIPS*, 2026 (**Spotlight Presentation**)

Automated Social Science: Language Models as Scientist and Subjects
(Benjamin S. Manning[1], **Kehang Zhu**[1], & John J. Horton)
Reject & Resubmit at *The Quarterly Journal of Economics*

Choose Your Agent: Tradeoffs in Adopting AI Advisors, Coaches, and Delegates in Multi-Party Negotiation
Collaboration with Google DeepMind
(**Kehang Zhu**[1],Nithum Thain Vivian Tsai, James Wexler, & Crystal Qian)
Under Review of CSCW

Engineering Simplicity for LLM Agents
(**Kehang Zhu**[1], Anand Shah[1],& David Parkes)
Under Review of *EC*

Benchmarking LLM Bidding Against Human Experimental Data
(Anand Shah[1], **Kehang Zhu**[1], Yancheng Jiang, Jeffrey Wang, Kerem Dayi, John Horton, & David Parkes)

[1]Indicates first author or co-first authors

Under Review of *EC*
Machine Behavior Workshop at *NeurIPS* , 2024

In Progress

An Experimental Study of AI Agent Delegation in multi-player Negotiation
(**Kehang Zhu**[1],Crystal Qian, John J. Horton, Benjamin S. Manning, Vivian Tsai, James Wexler, & Nithum Thain)
To be submitted to *Management Science*

Peer-Reviewed Publications

Strategic Tradeoffs Between Humans and AI in Multi-Agent Bargaining
Collaboration with Google DeepMind
(Crystal Qian[1], **Kehang Zhu**[1], John J. Horton, Benjamin S. Manning, Vivian Tsai, James Wexler, & Nithum Thain)
Proceedings of Conference of Intelligent User Interfaces (IUI), 2026

Reading between the pixels: Investigating the barriers to visualization literacy
(Carolina Nobre[1], **Kehang Zhu**, Eric Moerth, Hanspeter Pfister & Johanna Beyer)
Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems, 2024

Talks[2]

- **Conference on Artificial Intelligence, Machine Learning, and Business Analytics (2025)**
  Oral: An Experimental Study of AI Agent Delegation in multi-player Negotiation
- **Harvard EconCS Seminar (2025)**
  Oral: An Experimental Study of AI Agent Delegation in multi-player Negotiation
- **ACM Conference on Economics and Computation (EC 2025)**
  Poster: Strategic Tradeoffs Between Human and AI Agents in Bargaining Games.
- **NBER Summer Institute—Digital Economics and Artificial Intelligence Workshop (2024)**
  Plenary: Automated Social Science: Language Models as Scientist and Subjects
- **ACM Conference on Economics and Computation (EC 2024)**
  Workshop: Learning from Synthetic Labs: Language Models as Auction Participants.
- **International Conference on Computational Social Science (IC2S2 2024)**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **ACM Collective Intelligence Conference, Boston**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **Microsoft Research – AI, Cognition, and the Economy (AICE) Workshop**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **AI and the Future of Work Conference, Wharton School**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **Measuring Development, World Bank (MeasureDev 2024)**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **Econometric Society Interdisciplinary Frontiers Conference on Economics and AI+ML**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **American Causal Inference Conference (ACIC 2024)**
  Poster: Automated Social Science: Language Models as Scientist and Subjects
- **International Conference of the French Association of Experimental Economics (ASFEE 2024)**
  Oral: Automated Social Science: Language Models as Scientist and Subjects

[2]Includes co-author presentation

- **Data Science Summit (DataFun 2024)**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **AI Institute for Artificial Intelligence and Fundamental Interactions (2024)**
  Oral: Automated Social Science: Language Models as Scientist and Subjects
- **Interactive Causal Learning Conference (ICLC 2023)**
  Oral: Automated Social Science using Language Models and Structural Causal Models.
- **MIT Conference on Digital Experimentation (CODE 2023)**
  Oral: Automated Social Science using Language Models and Structural Causal Models

Teaching

15.S42: Leadership in the Age of Generative AI, **Guest Lecturer for John J. Horton**—MIT Sloan School of Management, Spring 2025

AI Summer School: Introduction to Language Models and Applications in Business; **Lecturer**, Harvard University, Summer 2024

AP-262: Statistical Physics; **TA for Vinothan N. Manoharan**—Harvard University, Fall 2023

Affiliations

Stanford Digital Economy Lab, Stanford Institute for Human-Centered Artificial Intelligence (HAI), MIT Initiative on the Digital Economy, Harvard EconCS Group

Review Service

Southern Economic Journal, Conference on Neural Information Processing Systems (NeurIPS), Conference on Digital Experimentation (CODE), ACM-Collective Intelligence(CI), International Conference on Computational Social Science (IC2S2), International Conference on Learning Representations (ICLR), International Conference on Machine Learning (ICML), ACM Conference on Human Factors in Computing Systems (CHI), ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW)

Academic Service

The Business Implications of Generative AI @ MIT (BIG.AI@MIT)

Employment

Amazon Science 2026
Applied Scientist Intern

Sloan School of Management, Massachusetts Institute of Technology 2023 - 2026
Research Assistant in Information Technology

Google DeepMind 2025
Research Intern

Harvard College 2024
AI Course Developer

Applify.ai 2022 - 2023
Natural Language Processing Researcher